

Transfer Learning

```
In [33]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import numpy as np
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import time
import os
import copy
import pretrainedmodels

torch.manual_seed(0) # for reproducibility
torch.cuda.empty_cache()
```

Load Data

We will use torchvision and torch.utils.data packages for loading the data.

```
In [34]: # Data augmentation and normalization for training
# Just normalization for validation
data_transforms = {
    'train': transforms.Compose([
        transforms.Resize([224, 224]),
        transforms.ToTensor()
    ]),
    'val': transforms.Compose([
        transforms.Resize([224, 224]),
        transforms.ToTensor()
    ]),
    'test': transforms.Compose([
        transforms.Resize([224, 224]),
        transforms.ToTensor()
    ])
}

data_dir = 'D:\data (augmented, 2 classes, tif)'
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
                                                data_transforms[x])
                  for x in ['train', 'val', 'test']}

batch_size = 128 # Need it as a global variable for computing average loss/accuracy per iteration
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=batch_size,
                                                shuffle=True, num_workers=4)
               for x in ['train', 'val', 'test']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val', 'test']}
class_names = image_datasets['train'].classes

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

Training the model

Now, let's write a general function to train a model. Here, we will illustrate:

- Scheduling the learning rate
- Saving the best model

In the following, parameter scheduler is an LR scheduler object from torch.optim.lr_scheduler.

```
In [35]: def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
    since = time.time()

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    epoch_numbers = []
    epoch_train_accuracies = []
    epoch_train_losses = []
    epoch_val_accuracies = []
    epoch_val_losses = []

    for epoch in range(num_epochs):
        epoch_numbers.append(epoch) # for plotting
        print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                if scheduler != None:
                    scheduler.step()
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

            running_loss = 0.0
            running_corrects = 0

            # Iterate over data.
            for inputs, labels in dataloaders[phase]: # The labels will correspond to the alphabetical order of the class names (https://discuss.pytorch.org/t/how-to-get-the-class-names-to-class-label-mapping/470):
                inputs = inputs.to(device)
                labels = labels.to(device)

                # zero the parameter gradients
                optimizer.zero_grad()

                # Forward
                # track history if only in train
                with torch.set_grad_enabled(phase == 'train'):
                    outputs = model(inputs)
                    _, preds = torch.max(outputs, 1)
                    loss = criterion(outputs, labels)

                # backward + optimize only if in training phase
                if phase == 'train':
                    loss.backward()
                    optimizer.step()

                # statistics
                running_loss += loss.item() * inputs.size(0)
                running_corrects += torch.sum(preds == labels.data)

            epoch_loss = running_loss/dataset_sizes[phase]
            epoch_acc = running_corrects.double()/dataset_sizes[phase]

            print('{} Loss: {:.4f} Acc: {:.4f}'.format(
                phase, epoch_loss, epoch_acc))

            # For plotting
            if phase == 'train':
                epoch_train_accuracies.append(epoch_acc)
                epoch_train_losses.append(epoch_loss)
            else:
                epoch_val_accuracies.append(epoch_acc)
                epoch_val_losses.append(epoch_loss)

            # deep copy the model
            if phase == 'val' and epoch_acc > best_acc:
                best_acc = epoch_acc
                best_model_wts = copy.deepcopy(model.state_dict())

        print()

        # Plotting

        plt.title("Training Curve (Loss)")
        plt.plot(epoch_numbers, epoch_train_losses, label="Train")
        plt.plot(epoch_numbers, epoch_val_losses, label="Validation")
        plt.xlabel("Epochs")
        plt.ylabel("Loss")
        plt.legend(loc='best')
        plt.show()

        plt.title("Training Curve (Accuracy)")
        plt.plot(epoch_numbers, epoch_train_accuracies, label="Train")
        plt.plot(epoch_numbers, epoch_val_accuracies, label="Validation")
        plt.xlabel("Epochs")
        plt.ylabel("Accuracy")
        plt.legend(loc='best')
        plt.show()

        time_elapsed = time.time() - since
        print('Training complete in {:.0fm} {:.0fs}'.format(
            time_elapsed // 60, time_elapsed % 60))
        print('Best val Acc: {:.4f}'.format(best_acc))

        # load best model weights
        model.load_state_dict(best_model_wts)
        return model
```

Finetuning the convnet

Load a pretrained model and reset final fully connected layer.

```
In [36]: model_ft = models.alexnet(pretrained=True)
model_ft.classifier[6] = nn.Linear(4096, len(class_names))
model_ft.classifier.add_module("7", nn.Dropout())

model_ft = model_ft.to(device)

criterion = nn.CrossEntropyLoss()

# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

optimizer_ft = optim.Adam([
    {'params': model_ft.features.parameters(), 'lr': 0.0001}, # The other (non-final) layers will have a lr = 0.1*base lr.
    {'params': model_ft.classifier[:6].parameters(), 'lr': 0.0001},
    {'params': model_ft.classifier[6:].parameters()} # The final layers will have the base lr.
], lr=0.001, weight_decay=0.005)
```

Train and evaluate

```
In [37]: model_ft = train_model(model_ft, criterion, optimizer_ft, scheduler="",
                                num_epochs=30)

Epoch 0/29
-----
train Loss: 0.5770 Acc: 0.6604
val Loss: 0.6674 Acc: 0.6226

Epoch 1/29
-----
train Loss: 0.5248 Acc: 0.7163
val Loss: 0.6427 Acc: 0.6827

Epoch 2/29
-----
train Loss: 0.4895 Acc: 0.7507
val Loss: 0.6153 Acc: 0.6724

Epoch 3/29
-----
train Loss: 0.4622 Acc: 0.7678
val Loss: 0.6815 Acc: 0.6947

Epoch 4/29
-----
train Loss: 0.4399 Acc: 0.7829
val Loss: 0.6598 Acc: 0.6587

Epoch 5/29
-----
train Loss: 0.4217 Acc: 0.7958
val Loss: 0.7217 Acc: 0.7033

Epoch 6/29
-----
train Loss: 0.4035 Acc: 0.8077
val Loss: 0.7331 Acc: 0.6861

Epoch 7/29
-----
train Loss: 0.3781 Acc: 0.8240
val Loss: 0.6605 Acc: 0.6895

Epoch 8/29
-----
train Loss: 0.3583 Acc: 0.8352
val Loss: 0.8524 Acc: 0.6792

Epoch 9/29
-----
train Loss: 0.3426 Acc: 0.8437
val Loss: 0.7691 Acc: 0.6930

Epoch 10/29
-----
train Loss: 0.3212 Acc: 0.8569
val Loss: 0.7768 Acc: 0.7084

Epoch 11/29
-----
train Loss: 0.3044 Acc: 0.8655
val Loss: 0.8321 Acc: 0.7033

Epoch 12/29
-----
train Loss: 0.2866 Acc: 0.8746
val Loss: 0.9156 Acc: 0.6810

Epoch 13/29
-----
train Loss: 0.2713 Acc: 0.8826
val Loss: 0.9528 Acc: 0.6741

Epoch 14/29
-----
train Loss: 0.2508 Acc: 0.8918
val Loss: 0.9184 Acc: 0.6930

Epoch 15/29
-----
train Loss: 0.2329 Acc: 0.9012
val Loss: 1.0556 Acc: 0.6913

Epoch 16/29
-----
train Loss: 0.2175 Acc: 0.9086
val Loss: 1.0164 Acc: 0.6844

Epoch 17/29
-----
train Loss: 0.2052 Acc: 0.9157
val Loss: 1.0020 Acc: 0.6535

Epoch 18/29
-----
train Loss: 0.1966 Acc: 0.9188
val Loss: 1.2152 Acc: 0.6861

Epoch 19/29
-----
train Loss: 0.1839 Acc: 0.9260
val Loss: 1.0866 Acc: 0.6621

Epoch 20/29
-----
train Loss: 0.1745 Acc: 0.9298
val Loss: 1.0988 Acc: 0.6861

Epoch 21/29
-----
train Loss: 0.1611 Acc: 0.9359
val Loss: 1.2685 Acc: 0.6792

Epoch 22/29
-----
train Loss: 0.1520 Acc: 0.9408
val Loss: 1.1721 Acc: 0.6792

Epoch 23/29
-----
train Loss: 0.1435 Acc: 0.9440
val Loss: 1.1314 Acc: 0.6758

Epoch 24/29
-----
train Loss: 0.1350 Acc: 0.9475
val Loss: 1.1743 Acc: 0.6655

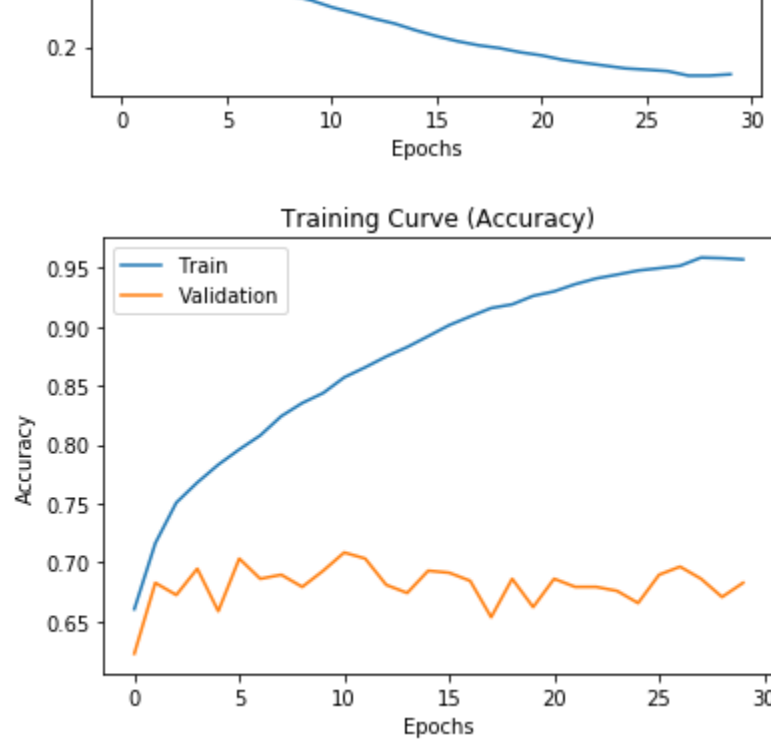
Epoch 25/29
-----
train Loss: 0.1311 Acc: 0.9495
val Loss: 1.2319 Acc: 0.6895

Epoch 26/29
-----
train Loss: 0.1267 Acc: 0.9515
val Loss: 1.0352 Acc: 0.6964

Epoch 27/29
-----
train Loss: 0.1129 Acc: 0.9585
val Loss: 1.3177 Acc: 0.6861

Epoch 28/29
-----
train Loss: 0.1130 Acc: 0.9580
val Loss: 1.0593 Acc: 0.6707

Epoch 29/29
-----
train Loss: 0.1171 Acc: 0.9569
val Loss: 1.1422 Acc: 0.6827
```



```
In [38]: dt = time.strftime("%Y%m%d-%H%M%S")
dt

Out[38]: '20190410-100203'
```

```
In [39]: # Save best model to disk for later!
torch.save(model_ft.state_dict(), os.getcwd() + '\\\\' + 'model_' + dt + '.pth')
```

```
In [40]: # Play sound when code finishes.
import winsound
duration = 1500 # milliseconds
```