

Transfer Learning

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import numpy as np
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import time
import os
import copy
import pretrainedmodels

torch.manual_seed(0) # for reproducibility
torch.cuda.empty_cache()
```

Load Data

We will use torchvision and torch.utils.data packages for loading the data.

```
In [2]: # Data augmentation and normalisation for training
# Just normalization for validation
data_transforms = {
    'train': transforms.Compose([
        transforms.Resize([224, 224]),
        transforms.ToTensor()
    ]),
    'val': transforms.Compose([
        transforms.Resize([224, 224]),
        transforms.ToTensor()
    ])
}

data_dir = 'D:\\data (augmented, 2 classes, tif)'
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
                                                    data_transforms[x])
                  for x in ['train', 'val']}

batch_size = 128 # Need it as a global variable for computing average loss/accuracy per iteration
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=batch_size,
                                                shuffle=True, num_workers=4)
               for x in ['train', 'val']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
class_names = image_datasets['train'].classes

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

Training the model

Now, let's write a general function to train a model. Here, we will illustrate:

- Scheduling the learning rate
- Saving the best model

In the following, parameter `scheduler` is an LR scheduler object from `torch.optim.lr_scheduler`.

```
In [3]: def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
    since = time.time()

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    epoch_numbers = []
    epoch_train_accuracies = []
    epoch_train_losses = []
    epoch_val_accuracies = []
    epoch_val_losses = []

    for epoch in range(num_epochs):
        epoch_numbers.append(epoch) # for plotting
        print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                if scheduler != '':
                    scheduler.step()
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

            running_loss = 0.0
            running_corrects = 0

            # Iterate over data.
            for inputs, labels in dataloaders[phase]: # The labels will correspond to the alphabetical order of the class names (https://discuss.pytorch.org/t/how-to-get-the-class-names-to-class-label-mapping/470).
                inputs = inputs.to(device)
                labels = labels.to(device)

                # zero the parameter gradients
                optimizer.zero_grad()

                # forward
                # track history if only in train
                with torch.set_grad_enabled(phase == 'train'):
                    outputs = model(inputs)
                    _, preds = torch.max(outputs, 1)
                    loss = criterion(outputs, labels)

                # backward + optimize only if in training phase
                if phase == 'train':
                    loss.backward()
                    optimizer.step()

                # statistics
                running_loss += loss.item() * inputs.size(0)
                running_corrects += torch.sum(preds == labels.data)

            epoch_loss = running_loss/dataset_sizes[phase]
            epoch_acc = running_corrects.double()/dataset_sizes[phase]

            print('{} Loss: {:.4f} Acc: {:.4f}'.format(
                phase, epoch_loss, epoch_acc))

            # For plotting
            if phase == 'train':
                epoch_train_accuracies.append(epoch_acc)
                epoch_train_losses.append(epoch_loss)
            else:
                epoch_val_accuracies.append(epoch_acc)
                epoch_val_losses.append(epoch_loss)

            # deep copy the model
            if phase == 'val' and epoch_acc > best_acc:
                best_acc = epoch_acc
                best_model_wts = copy.deepcopy(model.state_dict())

        print()

    # Plotting

    plt.title("Training Curve (Loss)")
    plt.plot(epoch_numbers, epoch_train_losses, label="Train")
    plt.plot(epoch_numbers, epoch_val_losses, label="Validation")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend(loc='best')
    plt.show()

    plt.title("Training Curve (Accuracy)")
    plt.plot(epoch_numbers, epoch_train_accuracies, label="Train")
    plt.plot(epoch_numbers, epoch_val_accuracies, label="Validation")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.legend(loc='best')
    plt.show()

    time_elapsed = time.time() - since
    print('Training complete in {:.0f}m {:.0f}s'.format(
        time_elapsed // 60, time_elapsed % 60))
    print('Best val Acc: {:.4f}'.format(best_acc))

    # load best model weights
    model.load_state_dict(best_model_wts)
    return model
```

Finetuning the convnet

Load a pretrained model and reset final fully connected layer.

```
In [4]: model_ft = models.resnet18(pretrained=True)
model_ft.fc = nn.Linear(512, len(class_names))

model_ft = model_ft.to(device)

criterion = nn.CrossEntropyLoss()

# All parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
```

Train and evaluate

```
In [5]: model_ft = train_model(model_ft, criterion, optimizer_ft, scheduler="",
                               num_epochs=30)

Epoch 0/29
-----
train Loss: 0.4962 Acc: 0.7442
val Loss: 0.6506 Acc: 0.6827

Epoch 1/29
-----
train Loss: 0.2642 Acc: 0.8912
val Loss: 0.8583 Acc: 0.6792

Epoch 2/29
-----
train Loss: 0.0927 Acc: 0.9699
val Loss: 1.1072 Acc: 0.7033

Epoch 3/29
-----
train Loss: 0.0316 Acc: 0.9919
val Loss: 1.3364 Acc: 0.6998

Epoch 4/29
-----
train Loss: 0.0111 Acc: 0.9982
val Loss: 1.4783 Acc: 0.6844

Epoch 5/29
-----
train Loss: 0.0101 Acc: 0.9976
val Loss: 1.6776 Acc: 0.6844

Epoch 6/29
-----
train Loss: 0.0034 Acc: 0.9997
val Loss: 1.7894 Acc: 0.6930

Epoch 7/29
-----
train Loss: 0.0086 Acc: 0.9980
val Loss: 1.8356 Acc: 0.6655

Epoch 8/29
-----
train Loss: 0.0035 Acc: 0.9995
val Loss: 1.8791 Acc: 0.6913

Epoch 9/29
-----
train Loss: 0.0019 Acc: 0.9998
val Loss: 2.0198 Acc: 0.6895

Epoch 10/29
-----
train Loss: 0.0061 Acc: 0.9985
val Loss: 1.8820 Acc: 0.6964

Epoch 11/29
-----
train Loss: 0.0028 Acc: 0.9993
val Loss: 2.0100 Acc: 0.6792

Epoch 12/29
-----
train Loss: 0.0013 Acc: 0.9999
val Loss: 2.0245 Acc: 0.6947

Epoch 13/29
-----
train Loss: 0.0011 Acc: 0.9999
val Loss: 2.0427 Acc: 0.6878

Epoch 14/29
-----
train Loss: 0.0013 Acc: 0.9997
val Loss: 2.0542 Acc: 0.7015

Epoch 15/29
-----
train Loss: 0.0007 Acc: 1.0000
val Loss: 2.0565 Acc: 0.6810

Epoch 16/29
-----
train Loss: 0.0006 Acc: 0.9999
val Loss: 2.0957 Acc: 0.6827

Epoch 17/29
-----
train Loss: 0.0012 Acc: 0.9998
val Loss: 2.1554 Acc: 0.6569

Epoch 18/29
-----
train Loss: 0.0059 Acc: 0.9983
val Loss: 2.1491 Acc: 0.6827

Epoch 19/29
-----
train Loss: 0.0012 Acc: 0.9998
val Loss: 2.0847 Acc: 0.6844

Epoch 20/29
-----
train Loss: 0.0011 Acc: 0.9998
val Loss: 2.1984 Acc: 0.6827

Epoch 21/29
-----
train Loss: 0.0005 Acc: 1.0000
val Loss: 2.2181 Acc: 0.6827

Epoch 22/29
-----
train Loss: 0.0004 Acc: 1.0000
val Loss: 2.2187 Acc: 0.6844

Epoch 23/29
-----
train Loss: 0.0005 Acc: 0.9999
val Loss: 2.2149 Acc: 0.6724

Epoch 24/29
-----
train Loss: 0.0004 Acc: 0.9999
val Loss: 2.2030 Acc: 0.6775

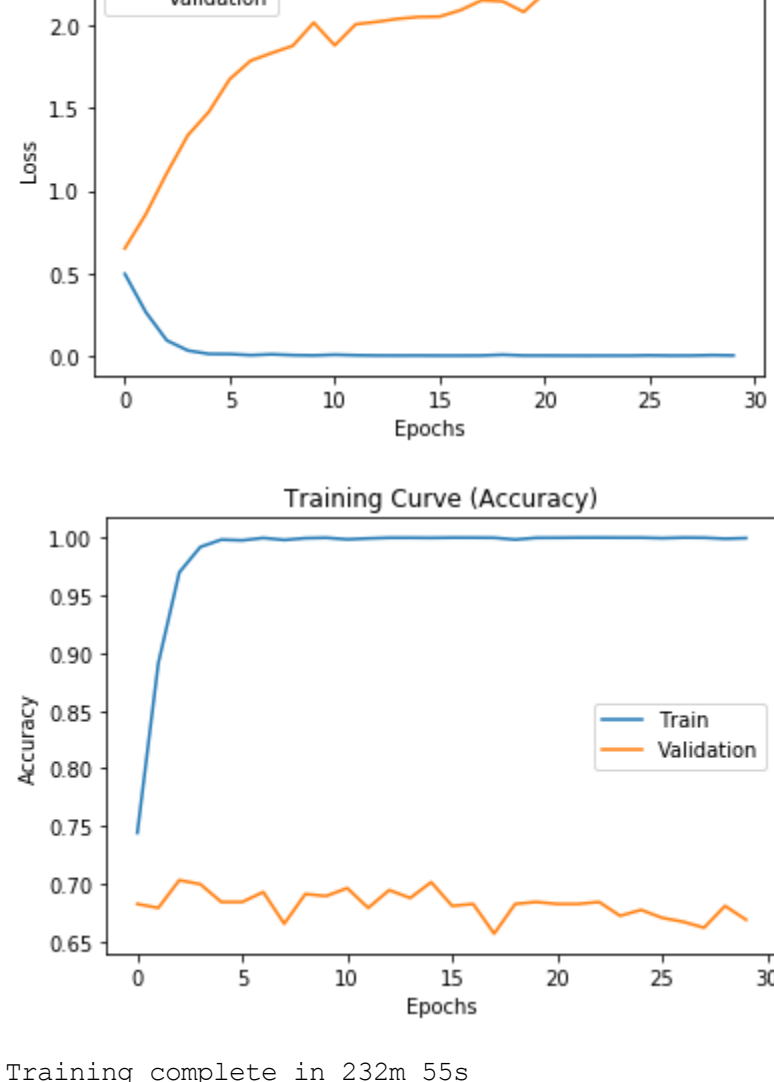
Epoch 25/29
-----
train Loss: 0.0022 Acc: 0.9995
val Loss: 2.2919 Acc: 0.6707

Epoch 26/29
-----
train Loss: 0.0006 Acc: 1.0000
val Loss: 2.3447 Acc: 0.6672

Epoch 27/29
-----
train Loss: 0.0009 Acc: 0.9999
val Loss: 2.4019 Acc: 0.6621

Epoch 28/29
-----
train Loss: 0.0034 Acc: 0.9990
val Loss: 2.3244 Acc: 0.6810

Epoch 29/29
-----
train Loss: 0.0016 Acc: 0.9995
val Loss: 2.3443 Acc: 0.6690
```



```
In [6]: dt = time.strftime("%Y%m%d-%H%M%S")
dt

Out[6]: '20190410-194241'
```

```
In [7]: # Save model to disk for later!
torch.save(model_ft.state_dict(), os.getcwd() + '\\\\' + 'model_' + dt + '.pth')
```

```
In [8]: # Play sound when code finishes.

import winsound
duration = 1500 # milliseconds
```